Meeting notes

2026-10-13

We discussed next directions to take after reading the Wysocki et al. (2022) article. We determined that the paper had done much of what we were planning to do (but, in some cases, in a cleaner way). We talked about how we'd have to position our study as an extension of that paper, and make sure that it was in fact contributing to the literature (and not simply replicating what they had done). We talked about the following potential ways that we could differentiate our paper / extend the Wysocki et al. paper:

- Explicitly testing the implication they make about data-driven approaches generally resulting in bias
- Examining Type I and Type II error rates (not just bias), which accounts for variance across research contexts and statistical significance
- Including two additional DGPs (irrelevant variables and independent causes of Y)
- Varying additional research contexts (e.g., amount of measurement error, relationship between C and Y)
- Examining cases where X has no effect on Y (given that false positives are often more problematic)

2026-10-6

We reviewed the PPT Elise had prepared for lab meeting, and discussed the following changes:

- Clarifying why covariates are used when analyzing experimental vs. observation data
- Moving the DGPs to the front (and referring to those terms, rather than "good" and "bad" covariates)
- Emphasizing the additional goal of our study (related to highlighting the problems with measurement error)
- Clarifying the steps to a simulation study
- Mentioning bias, not just Type I and Type II error rates

2026-09-29 & 2026-10-02

We executed our code though CHTC on a truncated version of jobs.csv, and successfully identified and fixed errors and warnings relating to string vs. numeric variables, where we sourced fun_cov2.R from, and namespace issues.

As next steps, Elise will:

- Prepare a 10-minute presentation for the Oct. 7 lab meeting, to frame a guided discussion for which research contexts to manipulate
- Remove calls to library
- Determine why we are getting 1000 jobs, instead of 500 [EG note: it's because we are running the code for two covariate selection approaches: all covariates, and no covariates]

Eventually, we can:

- Determine how many simulations we want to batch together
- Continue updating our combine.sh file to remove interim files, etc. (once we know everything is working)
- Decide which research contexts to manipulate (including DGPs), based on our discussion in lab meeting

2026-09-22

We reviewed and began updating the mak_jobs2.R, cov2.sub, and cov2.sh files. We also created error and results folders in our home directory.

As next steps, Elise will:

- Finish updating the three files to reflect the research contexts we are varying [EG note: I also need to update mak_jobs2.R to ensure that b_cy and b_cx appear as vectors in the format specified in fit_cov2.R]
- In mak_jobs2.R, define a more limited set of parameter values that we will vary to test out our code (including values that will require as much computational power as we expect we'll need, to be sure we have sufficient memory set up)
- Make sure she has access to the lab research drive

Next time, we will:

- Test out our code on a limited number of jobs
- Determine whether train.sif needs to be included in the list of transfer files in cov2.sub (by running a subfile that doesn't include it and seeing if we get an error)

Eventually, we may also need to:

- Update requested memory and disk parameters in cov2.sub
- Change file naming in the mak_jobs2.R file (if we want to run different batches and combine them)

2026-09-15

We began updating the fit_cov file to run simulations. We will start with the causal covariates framework, and vary parameters as follows:

- The effect of X on Y: 0 (no effects) and 0.3 (moderate effects)
- Sample size: 100, 200, 300, and 500
- Correlation between covariates and X: 0.1 (small effects) and 0.3 (moderate effects) [EG note: are these the path coefficients?]
- Correlation between covariates and Y: 0.1 (small effects) and 0.3 (moderate effects) [will be the same as above]
- Number of covariates: 1, 2, and 4
- Number of simulations per research setting: 1,000

As next steps, Elise will:

- Install WinSCP to transfer files
- Set empirical as a parameter (equal to TRUE or FALSE)
- Prepare to present at the Oct. 7 lab meeting to talk through which research contexts to vary
- Update parameter names in the fun_cov file so they match what we have in fit_cov
- Extract vectors from the bracketed format we'll use in our .csv file
- Generally tidy up the fit_cov file

Next time we meet, we will write CHTC code.

2026-09-04

To deal with the issue of means and SDs shifting from 0 and 1 (respectively) after we incorporate a causal structure, we determined that we will simply incorporate the causal relationships into our variance/covariance matrix (i.e., our sigma). We will specify the path diagrams and their coefficients, and calculate our sigma from there. In general, we will start with simple code, and build up from there to more complex scenarios.

We also talked about which parameters we care most about varying, since we can't vary them all. We decided that we care most about examining:

- The three causal structures (causal, distractor, and consequences of X)
- Zero and a non-zero effects of X on Y
- Reliability of covariates
- Sample size (N)

After that, we can consider:

- Number of covariates
- Size of the non-zero effect of X on Y
- Covariate effects on X and Y (where relevant) [EG note: this might matter if we use LASSO]
- Reliability of X
- Correlations among our covariates

We will return to this conversation at a later date, and consider factors like:

- What is going to have big effects
- When it will work vs. not work to use covariates
- Which scenarios people might not appreciate

As next steps, Elise will write code to create sigma from a path diagram, and confirm that this approach works (that is, that we get back out our path diagram coefficients after running regressions). We will start by using a very simple structure: the causal covariates context, with one covariate, and r_cov set to NULL. Elise will also begin thinking about what a two-covariate version (and eventually, an n-covariate version) would look like. (If that is proving too challenging, we can always just create separate statements for each number of covariates.) We will then run a simulation with the causal covariates context, and eventually loop back to the other contexts.

Finally, for future virtual meetings, we will record our meetings and use an agent to generate notes for us.

2026-08-21

We reviewed the function file from the first covariates study, and discussed creating functions for our three types of covariates (causal, distractor, and consequences of X). We decided that we would create X, Y, and covariates as normally-distributed continuous variables (mean = 0, SD = 1) that are uncorrelated with each other (in terms of the variance-covariance matrix), and then introduce causal relationships between them (based on the type of covariate).

2026-07-24

We decided that for now we will explore contexts where X has no effect on Y, *and* contexts where X has an effect on Y; this will allow us to explore methods for controlling for confounding variables while maintaining statistical power; if the story becomes too complicated, we can split our findings into two papers

We also discussed including three types of covariates, which we would keep in separate conditions (rather than mixing and matching them):

- Causal covariates
  - These cause both X and Y
  - These we would want to include in our models, because they are confounds
  - Example: when examining the relationship between parent-child relationship quality and child internalizing problems, we'd want to include child age, which influences both relationship quality and child internalizing problems
- Distractor covariates
  - These are not part of the model, and are uncorrelated with X, Y, and other covariates
  - These we wouldn't want to include in our models, because they'd reduce our power / make things inefficient
  - Example: when examining the relationship between intensive parenting beliefs and parental burnout, we would want to include parent gender (which can influence both intensive parenting beliefs and parental burnout), but not child gender (which likely doesn't influence intensive parenting beliefs or parental burnout and which is uncorrelated with parent gender)
- Consequences of X
  - These are caused by (not causes of) X, which are only correlated with Y because of their correlation with X
  - These we wouldn't want to include in our models, because we'd be underestimating the effect of X on Y
  - Example: when examining the relationship of child age on self-regulation, we wouldn't want to include child height, which is caused by child age but which doesn't itself impact self-regulation [EG Comment: I know this is a silly example, but I was having a lot of trouble coming up with something here that wasn't just a mediator!]

We may also consider a fourth type of covariate:

- Consequences of causal covariates
  - These are consequences of causal covariates, which are only correlated with Y because of their correlation with causal covariates
  - These we wouldn't want to include in our models, because they can decrease power / lead to inefficiencies, unless we use them in place of causal covariates
  - Example: when examining the relationship between parents' cultural values and child mental health, we'd want to include generational status but not primary language as a covariate, because generational status influences cultural values and child mental health; generational status also influences primary language, which does not independently influence child mental health

Finally, we decided that we will also likely include LASSO as a covariate selection approach

A revised summary of our study that reflects these updates is as follows: we will perform a simulation study investigating four common approaches to covariate selection across various research contexts, to determine which approaches best balance type I and type II errors (that is, that control for confounds while maintaining statistical power)

2026-07-16

We will perform a simulation study investigating three common approaches to covariate selection across various research contexts, to determine which approach results in the lowest average type I error

The covariate selection approaches are:

- Using all (theoretically-justified) covariates
- [JJC adds] Regressing Y on X and all covariates simultaneously
- Regressing Y on all covariates together (but not our focal variable) to determine which covariates to include
- Regressing Y on each covariate separately to determine which covariates to include

The research contexts we manipulate will be:

- Number of covariates
- Measurement error in our covariates
- Sample size
- Size of the effect of covariates on X
- Size of the effect of covariates on Y
- Correlations among covariates

We will test this using:

- A single focal variable X, with no true effect on Y, such that any effect we detect is a type I error; our focal variable will be continuous (and maybe eventually categorical)
- Continuous covariates only